

# Instacart Exploratory Data Analysis

## Project Introduction

In this exploratory data analysis, we are going to be analyzing Instacart customers in 2017 on their order preferences to better understand their shopping behavior.

## Analysis Outline

### Part A

- Verify that values in the 'order\_hour\_of\_day' columns in the orders table are in ranges from 0 to 23 and 'order\_dow' ranges from 0 to 6.
- Create a plot that shows how many people place orders for each hour of the day.
- Create a plot that shows what day of the week people shop for groceries.
- Create a plot that shows how long people wait until they place their next order.

### Part B

- Displayed the difference in 'order\_hour\_of\_day' distributions on Wednesdays and Saturdays using histograms.
- Plotted the distribution for the number of total orders that customers make.
- List the top 20 products that are most frequently ordered.

### Part C

- Calculated how large each customer's cart size is per order.
- List of the top 20 items that are reordered most frequently by our customers.

## Step 1: Openings Data Files

Importing libraries that will be used for the analysis

```
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
```

Importing instacart\_orders file

```
try:
    instacart_orders = pd.read_csv('instacart_orders.csv', sep=';')
except:
```

```
instacart_orders = pd.read_csv('/datasets/instacart_orders.csv',
sep=';')
```

Importing products file

```
try:
    products = pd.read_csv('products.csv', sep=';')
except:
    products = pd.read_csv('/datasets/products.csv', sep=';')
```

Importing order\_products file

```
try:
    order_products = pd.read_csv('order_products.csv', sep=';')
except:
    order_products = pd.read_csv('/datasets/order_products.csv',
sep=';')
```

Importing aisles file

```
try:
    aisles = pd.read_csv('aisles.csv', sep=';')
except:
    aisles = pd.read_csv('/datasets/aisles.csv', sep=';')
```

Importing departments file

```
try:
    departments = pd.read_csv('departments.csv', sep=';')
except:
    departments = pd.read_csv('/datasets/departments.csv', sep=';')
```

## Step 2: Pre-processing the Data

Verifying data for duplicates, null values, and incorrect data types

```
instacart_orders.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 478967 entries, 0 to 478966
Data columns (total 6 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   order_id            478967 non-null  int64
 1   user_id             478967 non-null  int64
 2   order_number        478967 non-null  int64
 3   order_dow           478967 non-null  int64
```

```

4    order_hour_of_day      478967 non-null  int64
5    days_since_prior_order  450148 non-null  float64
dtypes: float64(1), int64(5)
memory usage: 21.9 MB

```

In this DataFrame it has all the right data types and less than 10% missing values so we can say this DataFrame is set to be used

Now we are looking for duplicates and to remove them from the instacart\_orders dataset

```
instacart_orders.duplicated().sum()
```

```
15
```

```
instacart_orders[instacart_orders.duplicated()]
```

	order_id	user_id	order_number	order_dow	order_hour_of_day
\					
145574	794638	50898	24	3	2
223105	2160484	107525	16	3	2
230807	1918001	188546	14	3	2
266232	1782114	106752	1	3	2
273805	1112182	202304	84	3	2
284038	2845099	31189	11	3	2
311713	1021560	53767	3	3	2
321100	408114	68324	4	3	2
323900	1919531	191501	32	3	2
345917	2232988	82565	1	3	2
371905	391768	57671	19	3	2
394347	467134	63189	21	3	2
411408	1286742	183220	48	3	2
415163	2282673	86751	49	3	2
441599	2125197	14050	48	3	2
	days_since_prior_order				
145574	2.0				

223105	30.0
230807	16.0
266232	NaN
273805	6.0
284038	7.0
311713	9.0
321100	18.0
323900	7.0
345917	NaN
371905	10.0
394347	2.0
411408	4.0
415163	2.0
441599	3.0

Even though it is a small amount of duplicates compared to the total dataset it is still best to remove them

```
instacart_orders[(instacart_orders['order_dow'] == 4) &
(instacart_orders['order_hour_of_day'] == 2)].duplicated().sum()
0
```

Checking for duplicates in instacart\_orders DataFrame on Wednesdays at 2AM

```
instacart_orders.drop_duplicates(inplace=True)
```

Removing all duplicate orders from instacart\_orders DataFrame

```
instacart_orders.duplicated().sum()
0
```

Double checking for duplicates in the DataFrame

```
instacart_orders['order_id'].duplicated().sum()
0
```

Double checking for duplicates in the 'order\_id' column

Next, we are looking for duplicates and to remove them from the products dataset

All the data types look good and little data is missing

```
products.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 49694 entries, 0 to 49693
```

```
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   product_id       49694 non-null    int64
1   product_name      48436 non-null    object
2   aisle_id          49694 non-null    int64
3   department_id     49694 non-null    int64
dtypes: int64(3), object(1)
memory usage: 1.5+ MB

products.duplicated().sum()

0
```

No duplicates are found which is a good sign

```
products['product_id'].duplicated().sum()

0
```

Checking for duplicates in the 'product\_id' column and none were found

```
products['product_name'].duplicated().sum()

1257
```

After checking for duplicates in the 'product\_name' column we have a noticeable amount of duplicates. Since this column is an object type we can check to see if these are actual duplicates or a character casing issue.

```
products['product_name'].str.lower()

0           chocolate sandwich cookies
1           all-seasons salt
2           robust golden unsweetened oolong tea
3   smart ones classic favorites mini rigatoni wit...
4           green chile anytime sauce
...
49689          high performance energy drink
49690          original pancake & waffle mix
49691   organic instant oatmeal light maple brown sugar
49692          spring water body wash
49693          burrito- steak & cheese
Name: product_name, Length: 49694, dtype: object
```

By using str.lower() method we turned all the string in this column to lowercase

```
products['product_name'].duplicated().sum()
```

1257

```
products[products['product_name'].notna()  
['product_name'].str.lower().duplicated().sum()]
```

104

After lowering the case of all entries in that column we have an increase in duplicates

```
products[~products['product_name'].duplicated()]
```

	product_id	
product_name \		
0	1	Chocolate Sandwich Cookies
1	2	All-Seasons Salt
2	3	Robust Golden Unsweetened Oolong Tea
3	4	Smart Ones Classic Favorites Mini Rigatoni Wit...
4	5	Green Chile Anytime Sauce
...	...	...
49689	49690	HIGH PERFORMANCE ENERGY DRINK
49690	49691	ORIGINAL PANCAKE & WAFFLE MIX
49691	49692	ORGANIC INSTANT OATMEAL LIGHT MAPLE BROWN SUGAR
49692	49693	SPRING WATER BODY WASH
49693	49694	BURRITO- STEAK & CHEESE

	aisle_id	department_id
0	61	19
1	104	13
2	94	7
3	38	1
4	5	13
...	...	...
49689	64	7
49690	130	14
49691	130	14
49692	127	11
49693	38	1

[48437 rows x 4 columns]

Checking for duplicate product names that aren't missing

Now we move on to the next dataset to search and remove duplicates from the order\_products dataset. Since this dataset is so large we need to use the "show\_counts=True" parameter to show all the rows

```
order_products.info(show_counts=True)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4545007 entries, 0 to 4545006
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   order_id              4545007 non-null  int64
1   product_id           4545007 non-null  int64
2   add_to_cart_order     4544171 non-null  float64
3   reordered             4545007 non-null  int64
dtypes: float64(1), int64(3)
memory usage: 138.7 MB

order_products.duplicated().sum()

0
```

Next, we move onto the aisles dataset to search for and remove all duplicates if possible

```
aisles.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 134 entries, 0 to 133
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   aisle_id    134 non-null    int64
1   aisle       134 non-null    object
dtypes: int64(1), object(1)
memory usage: 2.2+ KB

aisles.duplicated().sum()

0
```

Lastly, we need to clean the data for the departments dataset

```
departments.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
```

```

0    department_id    21 non-null    int64
1    department      21 non-null    object
dtypes: int64(1), object(1)
memory usage: 468.0+ bytes

departments.duplicated().sum()

0

```

## Exploring Missing Values

```

products.isna().sum()

product_id          0
product_name      1258
aisle_id           0
department_id       0
dtype: int64

```

Checking for missing values and one is found in the 'product\_name' column

```

print(products[products['product_name'].isna()])

   product_id product_name aisle_id department_id
37          38         NaN        100           21
71          72         NaN        100           21
109         110         NaN        100           21
296         297         NaN        100           21
416         417         NaN        100           21
...         ...         ...         ...         ...
49552       49553         NaN        100           21
49574       49575         NaN        100           21
49640       49641         NaN        100           21
49663       49664         NaN        100           21
49668       49669         NaN        100           21

[1258 rows x 4 columns]

```

By filtering the 'product\_name' column we found the one column with missing values

```

print(products[products['aisle_id'] == 100]['product_name'].isna())

37      True
71      True
109     True
296     True
416     True
...
49552   True
49574   True

```



```
49640    True
49663    True
49668    True
Name: product_name, Length: 1258, dtype: bool
```

By looking more into 'aisle\_id' column it shows that the aisle is missing the same amount of values as the 'product\_name' column. We can conclude that all the missing 'product\_name' values are the same ones as in the 'aisle\_id' 100.

```
print(products[products['department_id'] == 21]
      ['product_name'].isna())

37      True
71      True
109     True
296     True
416     True
...
49552   True
49574   True
49640   True
49663   True
49668   True
Name: product_name, Length: 1258, dtype: bool
```

By looking more into 'department\_id' column it shows that the department is missing the same amount of values as the 'product\_name' column and the 'aisle\_id' column. We can conclude that all the missing 'product\_name' values are the same ones as in the 'department\_id' 21.

```
products['product_name'] = products['product_name'].fillna('unknown')
```

Renaming all the NaN values to 'unknown'

```
products.isna().sum()

product_id      0
product_name    0
aisle_id        0
department_id   0
dtype: int64

products[products['product_name'].isna()]

Empty DataFrame
Columns: [product_id, product_name, aisle_id, department_id]
Index: []

products[products['product_name'] == 'unknown']
```

	product_id	product_name	aisle_id	department_id
37	38	unknown	100	21
71	72	unknown	100	21
109	110	unknown	100	21
296	297	unknown	100	21
416	417	unknown	100	21
...	...	...	...	...
49552	49553	unknown	100	21
49574	49575	unknown	100	21
49640	49641	unknown	100	21
49663	49664	unknown	100	21
49668	49669	unknown	100	21

[1258 rows x 4 columns]

Double checking for null values

```
instacart_orders.isna().sum()
```

```
order_id          0
user_id           0
order_number      0
order_dow         0
order_hour_of_day 0
days_since_prior_order    28817
dtype: int64
```

Checking for missing values in the instacart\_orders dataframe

```
print(instacart_orders[(instacart_orders['days_since_prior_order'].isna() & (instacart_orders['order_number'] > 1))])
```

Empty DataFrame

```
Columns: [order_id, user_id, order_number, order_dow,
order_hour_of_day, days_since_prior_order]
Index: []
```

After filtering the null values and those who are not ordering for the first time it resulted in an empty dataframe meaning that there are no null values in the days\_since\_prior\_order column that are not first time customers

```
print(order_products.isna().sum())
```

```
order_id          0
product_id        0
add_to_cart_order    836
reordered          0
dtype: int64
```

```
order_products[order_products['add_to_cart_order'].isna()]
```

	order_id	product_id	add_to_cart_order	reordered
737	2449164	5068	NaN	0
9926	1968313	43867	NaN	0
14394	2926893	11688	NaN	0
16418	1717990	4142	NaN	0
30114	1959075	42828	NaN	1
...	...	...	...	...
4505662	1800005	7411	NaN	0
4511400	1633337	260	NaN	0
4517562	404157	9517	NaN	0
4534112	1673227	17835	NaN	0
4535739	1832957	17949	NaN	1

```
[836 rows x 4 columns]
```

After looking for missing values in the 'order\_products' DataFrame there are a lot of missing orders in the 'add\_to\_cart\_order' column

```
print(order_products['add_to_cart_order'].min())
print(order_products['add_to_cart_order'].max())
```

```
1.0
64.0
```

After printing the minimum and maximum values from the 'add\_to\_cart\_order' column it shows that the min value is 1 and the max value is 64

```
null_order_id =
order_products[order_products['add_to_cart_order'].isna()]
['order_id'].to_list()
```

Saved all order IDs with at least one missing value in 'add\_to\_cart\_order' to null\_order\_id

```
orders_missing_cart_order =
order_products[order_products['order_id'].isin(null_order_id)]
```

After saving orderIDs to null\_order\_id, we filtered the 'order\_id' column in the order\_products DataFrame with the new null\_order\_id list and saved the results to orders\_missing\_cart\_order variable

```
orders_missing_cart_order.groupby('order_id')
['order_id'].count().min()
```

```
65
```

```
orders_missing_cart_order.groupby('order_id').size().min()
```

65

When grouping the dataframe with the users with missing add\_to\_cart\_orders, it shows that the minimum value of their missing orders starts at 65. We know beforehand that the max value of add\_to\_cart\_orders was 64, therefore we can conclude that there is a technical issue with the recording method after the add\_to\_cart\_order reaches 64.

```
order_products['add_to_cart_order'] =  
order_products['add_to_cart_order'].fillna('999')
```

Renamed all null rows to 999 to remove all the null values

```
print(order_products[order_products['add_to_cart_order'].isna()])  
  
Empty DataFrame  
Columns: [order_id, product_id, add_to_cart_order, reordered]  
Index: []
```

Double checking for null values

```
order_products['add_to_cart_order'] =  
order_products['add_to_cart_order'].astype('int')
```

Changing the column type to integer type

```
order_products.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4545007 entries, 0 to 4545006  
Data columns (total 4 columns):  
#   Column                Dtype  
---  -  
0   order_id              int64  
1   product_id            int64  
2   add_to_cart_order     int32  
3   reordered             int64  
dtypes: int32(1), int64(3)  
memory usage: 121.4 MB
```

Checking to see if the column data type has changed to an integer type

## Step 3: Data analysis

### Part: A

Part A-1: To verify 'order\_hour\_of\_day' values are 0-23 and 'order\_dow' are 0-6

To do this I will be using the min and max values on both of these columns to see if they meet the requirements

```
print(instacart_orders['order_hour_of_day'].min())
print(instacart_orders['order_hour_of_day'].max())

0
23

print(instacart_orders['order_dow'].min())
print(instacart_orders['order_dow'].max())

0
6
```

Since the minimum and maximum values are within the value ranges it means that all the values are accurate and there are no outlier data points that need to be fixed

Part A-2: Create a plot that shows how many people placed orders for each hour of the day


First I am going to group the instacart\_orders by 'order\_hour\_of\_day' to separate the data results into each hour. Then I will count 'order\_id' to show each transaction for the grouped hours.

```
hourly_orders = instacart_orders.groupby('order_hour_of_day')
['order_id'].count()
```

To create this plot I will need to use matplotlib

```
hourly_orders.plot(kind='bar',
                    x='order_hour_of_day',
                    y='order_id',
                    title='Total Orders per Hour',
                    xlabel='Hour of the Day',
                    ylabel='Total Orders')

plt.grid(axis='y')
plt.show()
```



In this graph it shows that orders tend to start trending up at 7am generally when customers start to wake up and start to trend downward around 4pm. In this dataset it shows that Instacart customers tend to order the most during the 9-5 work day.


Part A-3: Create a plot that shows what day of the week people shop for groceries

First I am going to group the instacart\_orders by 'order\_dow' to separate the data into each day of the week. Then I will count 'order\_id' to show each transaction for the grouped day.

```
dow_orders = instacart_orders.groupby('order_dow')['order_id'].count()
```

To create this plot I will need to use matplotlib

```
dow_orders.plot(kind='bar',
                 x='order_dow',
                 y='order_id',
                 title='Day of the Week Total Orders',
                 xlabel='Day of the Week',
                 ylabel='Total Orders')
plt.grid(axis='y')
plt.show()
```



On the day of the week axis it would be safe to assume that 0 is correlated to Saturday. Saturdays are usually the busiest days for most businesses since customers have time off to run their errands like grocery shopping and this is also supported as 4 would be correlated to Wednesday as Wednesday is usually the slowest day of the week for most businesses.

In this graphic it shows that the first two days of the week (Saturday and Sunday) are the most busy and that Instacart should allocate more resources into those days as those are the times customers are more likely to stop by and that day 4 (Wednesday) is the slowest day of the week.

Part A-4: Create a plot that shows how long people wait until placing their next order

```
print(instacart_orders['days_since_prior_order'].mean())
print(instacart_orders['days_since_prior_order'].median())
print(instacart_orders['days_since_prior_order'].min())
print(instacart_orders['days_since_prior_order'].max())
```

```
11.101813900274362
7.0
0.0
30.0
```

From the mean, median, min, and max methods it shows that the median customer will wait about a week before they reorder. Also, I should note that since the average is significantly higher than the median it shows that there are some large outliers when it comes to small parts of the customers.

```
days_unitl_reorder =  
instacart_orders['days_since_prior_order'].value_counts()
```

By using value\_counts() it will separate the frequency of which days the customers will reorder

```
print(days_unitl_reorder.head(10))
```

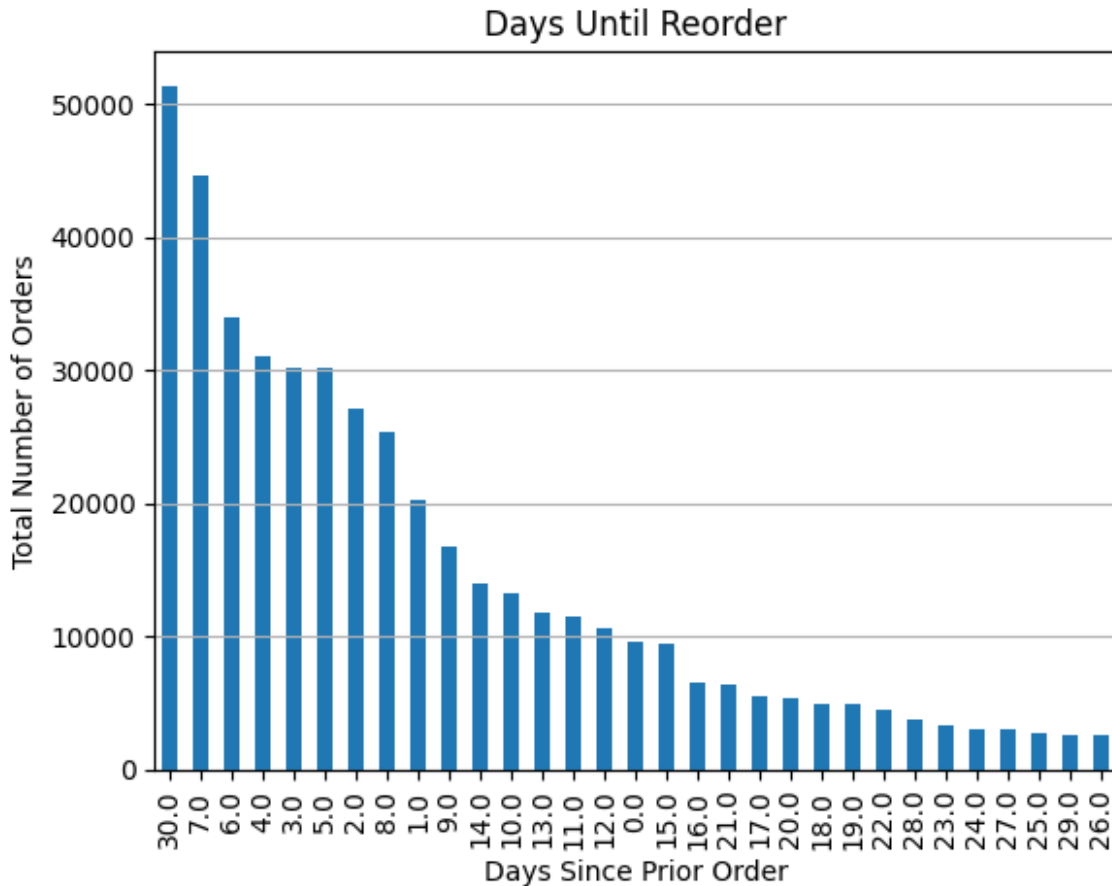
```
days_since_prior_order  
30.0    51337  
7.0     44577  
6.0     33930  
4.0     31006  
3.0     30224  
5.0     30096  
2.0     27138  
8.0     25361  
1.0     20179  
9.0     16753  
Name: count, dtype: int64
```

In this dataset it shows that most popular time frames that customers order is at least once a week or exactly once a month.

This makes sense as when we compared the mean to the median that about half order at least once a week and that there is a large outlier of customers that take a lot longer to order skewing the average up from the median value.

```
days_unitl_reorder.plot(kind='bar',  
                          title='Days Until Reorder',  
                          xlabel='Days Since Prior Order',  
                          ylabel='Total Number of Orders')  
plt.grid(axis='y')  
plt.show()
```





In this graphic it shows that a heavy majority of customers either reorder on Instacart within 7 days or they order exactly once a month

## Part: B

Part B-1: Finding the difference in hourly order sales volume between Wednesdays and Saturdays

```
wednesday = instacart_orders[instacart_orders['order_dow'] == 4]
```


First, I need to isolate the orders for just Wednesday

```
wed_hourly_orders = wednesday.groupby('order_hour_of_day')
['order_id'].count()
```

Now we use the new wednesday variable and count the hourly orders

```
wed_hourly_orders.plot(kind='bar',
                        x='order_hour_of_day',
                        y='order_id',
                        title='Wednesday Hourly Orders',
                        xlabel='Hour of the Day',
                        ylabel='Total Orders')
```

```
plt.grid(axis='y')
plt.show()
```



The distribution in orders on Wednesday looks very similar to the average overall with two peaks during the busiest hours

```
saturday = instacart_orders[instacart_orders['order_dow'] == 0]
```


Same as Wednesday I need to isolate the orders for just Saturday

```
sat_hourly_orders = saturday.groupby('order_hour_of_day')
['order_id'].count()
```

After finding all the orders for Saturday we group the orders by the hour and count the number of orders in each hour

```
sat_hourly_orders.plot(kind='bar',
                        x='order_hour_of_day',
                        y='order_id',
                        title='Saturday Hourly Orders',
                        xlabel='Hour of the Day',
                        ylabel='Total Orders')
```

```
plt.grid(axis='y')
plt.show()
```



Unlike the graph for Wednesday, the Saturday graph has one peak in orders instead of two peaks in the Wednesday graph. Also the volume of orders is noticeably higher throughout the day

```
instacart_orders[instacart_orders['order_dow'] == 0]
['order_hour_of_day'].hist(alpha=.5,

bins=24,


color='blue')
instacart_orders[instacart_orders['order_dow'] == 4]
['order_hour_of_day'].hist(alpha=.5,

bins=24,

color='red')

plt.xlabel('Hour of the Day')
plt.ylabel('Total Orders')
```

```
plt.title('Saturday VS Wednesday Hourly Orders')
plt.show()
```



When we overlay the two graphs for Saturday and Wednesday it looks like Saturday clearly does more volume in orders compared to Wednesday. This difference can be seen especially towards the peak hours of the day.


Part B-2: Customer order distribution plot

```
customer_order_freq = instacart_orders.groupby('user_id')
['order_number'].count()
```

First, I grouped all the customers by 'user\_id' and added all their 'order\_numbers' and saved it under the variable 'customer\_order\_freq'

```
customer_order_freq.plot(kind='hist',
                        title='Customer Order Frequency',
                        xlabel='Orders Placed',
                        ylabel='Order Frequency')

plt.grid(axis='y')
plt.show()
```



In this graphic it shows that most customers place less than 25 orders

Part B-3: Top 20 Products

```
product_merged = order_products.merge(products,  
                                       on='product_id')
```

First, I need to merge two DataFrames that contained the same product ids to get the product name

```
product_merged.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4545007 entries, 0 to 4545006  
Data columns (total 7 columns):  
#   Column                Dtype  
---  ---  
0   order_id              int64  
1   product_id            int64  
2   add_to_cart_order     int32  
3   reordered             int64  
4   product_name          object  
5   aisle_id              int64  
6   department_id         int64
```


```
dtypes: int32(1), int64(5), object(1)
memory usage: 225.4+ MB
```

I used info to check if the merge changed any of the data types and as expected nothing changed

```
top_20_list = product_merged.groupby(['product_name', 'product_id'])
['order_id'].count().sort_values(ascending=False).head(20)
```

To find the top 20 products I grouped the products by 'product\_name' and 'product\_id', counted each order in those groups, and sorted the values from highest to lowest. By getting the first 20 values will be the top 20 products since it is sorted from largest value to lowest value.

```
top_20_list.plot(kind='bar',
                  title='Top 20 Sold Products by Volume',
                  xlabel='Top 20 Product Names and IDs',
                  ylabel= 'Sales Volume')
plt.grid(axis='y')
plt.show()
```



Based on the top 20 products sold a majority of the volume sold seems to be fruits and vegetables

## Part: C

Part C-1: Cart size per order

First, I have to merge the DataFrames I will need on this plot. I will merge the DataFrames on the 'order\_id' column as that is the column that both DataFrames have in common.

```
cart_size_merged = order_products.merge(instacart_orders,  
                                         on='order_id')
```

Then I need to check if all the columns I need will be there and that there is no change in their data type

```
cart_size_merged.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4545007 entries, 0 to 4545006  
Data columns (total 9 columns):  
#   Column                Dtype  
---  -----  
0   order_id              int64  
1   product_id            int64  
2   add_to_cart_order     int32  
3   reordered             int64  
4   user_id               int64  
5   order_number          int64  
6   order_dow             int64  
7   order_hour_of_day     int64  
8   days_since_prior_order float64  
dtypes: float64(1), int32(1), int64(7)  
memory usage: 294.7 MB
```

I created the 'customer\_cart\_size' variable to group the the dataset by order and how many products are in each order

```
customer_cart_size = cart_size_merged.groupby('order_id')  
['product_id'].count()  
  
customer_cart_size.describe()  
  
count      450046.000000  
mean         10.098983  
std           7.540206  
min           1.000000  
25%           5.000000  
50%           8.000000  
75%          14.000000  
max          127.000000  
Name: product_id, dtype: float64
```

Looking more into the customer\_cart\_size, it shows that the average cart size is about 10 items, and the median cart size is 8 items. Since the average is noticeably larger than the median it shows that there are some large outliers cart sizes increasing the average. Looking at the bell chart it shows that the max cart size is 127 orders.


Now we can find the frequency of each cart size per order and graph the top 20 most popular cart sizes



```
customer_cart_size = customer_cart_size.value_counts().head(20)
```

After finding the frequency of each order size we can now graph the results to see the most popular cart size

```
customer_cart_size.plot(kind='bar',  
                        title='Top 20 Customer Order Size Frequency',  
                        xlabel='Cart Size',  
                        ylabel='Cart Size Frequency')  
  
plt.grid(axis='y')  
plt.show()
```



Based on the plot, it displays that the most popular cart sizes are 5, 6, 4, 7, and 3.


Part C-2: What are the top 20 items that are reordered most frequently?

```
top_20_reorders = product_merged.groupby(['product_name',  
                                           'product_id'])  
['reordered'].sum().sort_values(ascending=False).head(20)
```

Firstly, I grouped the products by name and id from the custom DataFrame 'product\_merged' we created earlier. Then I added all the reorder values in each order together. Since reordered items will show as a 1 if it is reordered and 0 if it is not then we can tally up the about of reorders

per product and get the most popular reordered item from there. Lastly, I sorted the values from largest to smallest and saved the top 20 values. Since the items are sorted from highest to lowest the first 20 items will be the 20 most popular reordered items.

```
top_20_reorders.plot(kind='bar',  
                      title='Top 20 Reordered Items',  
                      xlabel='Top 20 Reordered Items and Item ID',  
                      ylabel='Frequency of Purchase')  
plt.grid(axis='y')  
plt.show()
```



Many of the items on this list are also items on the top 20 sold items as well showing that the top 20 sold items are also the most repurchased items

## Conclusion

This analysis of Instacart's 2017 data shows that most customers are weekly shoppers or once a month shoppers. These customers often come to buy fruits and vegetables indicating that the produce Instacart sells are possibly of high quality and could possibly sold for a higher premium. Instacart customers also tend to order more on the weekends and show that there could be more room for improvement on the weekdays. Possibly, if there were more sales during the slower days like Wednesday it could attract more customers and have them be more frequent customers.